InfluxDB vs. APM and logging apps

We often get asked how InfluxData is different from applications like Datadog, SumoLogic, Splunk, New Relic, and other metrics and monitoring systems. First and foremost, InfluxData is purpose-built as a time series data platform. It exists so that developers can build their applications on top of the platform.

The above-mentioned applications have their user interfaces and business logic built-in. They're meant to be off-the-shelf solutions that give developers whatever they need for the given problem. Most developers wouldn't build their custom applications on top of these bundled solutions.

InfluxData is a platform for developers to build upon. It's meant to be customized for the unique business logic of each organization it's deployed by. This makes it an ideal choice for larger organizations that are looking to develop solutions designed specifically for their needs or for application developers creating solutions for customer-facing products.

InfluxDB vs. other databases

InfluxDB is often compared to other databases. However, when doing a comparison, the entirety of the InfluxDB platform should be taken into account. There are multiple types of databases that get brought up for comparison. Mostly these are distributed databases like Cassandra or more time-series-focused databases like Graphite or RRD.

When comparing InfluxData with Cassandra or HBase, there are some stark differences. First, those databases require a significant investment in developer time and code to recreate the functionality provided out of the box by InfluxDB. Specifically, developers will need to write code to shard the data across the cluster, aggregate and downsample functions, data eviction and lifecycle management, and summarization. Finally, they'll have to create an API to write and query their new service.

When the rest of the InfluxDB platform is brought into the picture, developers using Cassandra or HBase have even more ground to make up. They'll need to write tools for data collection, introduce a real-time processing system and write code for monitoring and alerting. Finally, they'll need to write a visualization engine to display the time series data to the user. While some of these tasks are handled with other time series databases, there are a few key differences between the other solutions and InfluxDB. First, other time series solutions like Graphite or OpenTSDB are designed with only regular time series data in mind and don't have the ability to store raw high-precision data and downsample it on the fly.

While with other time series databases, the developer must summarize their data before they put it into the database, InfluxDB lets the developer seamlessly transition from raw time series data into summarizations.