InfluxDB vs. stream processing solutions

Stream processing solutions like Spark, Storm, Kafka and others are often used for real-time analytics and other time series processing workloads. With InfluxData's Kapacitor, there is some overlap with these solutions. However, these are mostly complementary tools that users have already included in their InfluxDB architectures.

With that being said, we can look at how Kapacitor differs from these other general-purpose stream processing tools. First, Kapacitor has the structure and schema of InfluxDB as a core part of its processing engine. It has a domain-specific language built in to help developers do common tasks on InfluxDB time series.

For developers using a general-purpose processing engine, they'll need to write code to decode, process, and re-encode time series data. Kapacitor comes with these tools out of the box letting the developer think more about what they want to monitor and less about the tooling needed to get the job done.

Additionally, the development of Flux, InfluxData's new scripting and query language, has made InfluxDB's time series analytics capabilities all the more powerful.

InfluxData *is* time series data

Fueled by the massive growth of connected devices (i.e., IoT) and rapidly increasing instrumentation requirements of next-generation software, time series technology has become more popular. Since launching InfluxDB, an open source time series platform in 2013, we have seen millions of downloads, built an expanding list of enterprise customers, and fostered a growing community that is always finding new ways to deploy and build on our platform. InfluxData also offers two paid editions of InfluxDB: InfluxDB Cloud (managed database as a service) and InfluxDB Enterprise (subscription that turns any InfluxData instance into a production-ready cluster that can run anywhere).

InfluxData has a narrow focus for what has rapidly become a horizontal use case. Our narrow focus means that the entire stack can have optimizations for performance and developer productivity that other general-purpose solutions can't match, such as high compression, super-fast engines and powerful query language designed to best work with flow-based models. At the same time, InfluxData provides a platform that is broadly customizable, making it a perfect choice for developers who want greater control over their tooling than what out-of-the-box applications and solutions provide.